

AI & ML Internship Task-3

By Elevate Labs

1.Import and preprocess the dataset.

```
In [24]: import pandas as pd

# Load your horse dataset
df = pd.read_csv('Housing.csv')

# Replace 'yes' with 1 and 'no' with 0
df.replace({'yes': 1, 'no': 0}, inplace=True)

df.replace({'furnished' :1,'unfurnished':2,'semi-furnished':3},inplace=True)

# (No saving to CSV)
print(df.head())
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	
2	12250000	9960	3	2	2	1	0	
3	12215000	7500	4	2	2	1	0	
4	11410000	7420	4	1	2	1	1	

	basement	hotwaterheating	airconditioning	parking	prefarea	\
0	0	0	1	2	1	
1	0	0	1	3	0	
2	1	0	0	2	1	
3	1	0	1	3	1	
4	1	0	1	2	0	

	furnishingstatus
0	1
1	1
2	3
3	1
4	1

C:\Users\sujal\AppData\Local\Temp\ipykernel_13972\2031134496.py:7: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call 'result.infer_object_s(copy=False)'. To opt-in to the future behavior, set 'pd.set_option('future.no_silent_downcasting', True)'
df.replace({'yes': 1, 'no': 0}, inplace=True)
C:\Users\sujal\AppData\Local\Temp\ipykernel_13972\2031134496.py:9: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call 'result.infer_object_s(copy=False)'. To opt-in to the future behavior, set 'pd.set_option('future.no_silent_downcasting', True)'
df.replace({'furnished' :1,'unfurnished':2,'semi-furnished':3},inplace=True)

```
In [26]: df
```

Out [26]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	13300000	7420	4	2	3	1	0	0	0	1	2	1	1
1	12250000	8960	4	4	4	1	0	0	0	1	3	0	1
2	12250000	9960	3	2	2	1	0	1	0	0	2	1	3
3	12215000	7500	4	2	2	1	0	1	0	1	3	1	1
4	11410000	7420	4	1	2	1	1	1	0	1	2	0	1
...
540	1820000	3000	2	1	1	1	0	1	0	0	2	0	2
541	1767150	2400	3	1	1	0	0	0	0	0	0	0	3
542	1750000	3620	2	1	1	1	0	0	0	0	0	0	2
543	1750000	2910	3	1	1	0	0	0	0	0	0	0	1
544	1750000	3850	3	1	2	1	0	0	0	0	0	0	2

545 rows x 13 columns

2.Split data into train-test sets.

```
In [57]: X = df.iloc[:,1:13]
y = df['price']

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score,mean_absolute_error

X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.2, random_state=50)
```

3.Fit a Linear Regression model using sklearn.linear_model.

4.Evaluate model using MAE, MSE, R².

```
In [59]: lr_model = LinearRegression()
lr_model.fit(X_train,y_train)

y_pred = lr_model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)


print("Mean Squared Error:", mse)
print("R2 Score:", r2)
print("mean absolute error is :", mae)
```

Mean Squared Error: 852563443450.0237
R2 Score: 0.749309943628877
mean absolute error is : 745316.921824061

5.Plot regression line and interpret coefficients.

```
In [66]: import matplotlib.pyplot as plt

plt.scatter(y_test, y_pred, color='blue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linewidth=2) # perfect line
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Actual vs Predicted Price')
plt.grid(True)
plt.show()
```



```
In [64]: print("Intercept:", lr_model.intercept_)
coefficients = pd.DataFrame({'Feature': X.columns, 'Coefficient': lr_model.coef_})
print(coefficients)
```

Intercept: -407110.98745003715

	Feature	Coefficient
0	area	238.037531
1	bedrooms	129501.457438
2	bathrooms	949101.720050
3	stories	476971.998073
4	mainroad	584161.590414
5	guestroom	264903.335677
6	basement	332312.325001
7	hotwaterheating	911223.087147
8	airconditioning	916475.605650
9	parking	323939.763316

10	prefarea	685541.256199
11	furnishingstatus	51808.496467
In []:		